

Computation of Standardized Residuals for MARSS Models

E. E. Holmes*

January 3, 2020

Abstract

This report shows how to compute the variance of the joint conditional model and state residuals for multivariate autoregressive Gaussian state-space (MARSS) models. The bulk of the report focuses on ‘smoothations’, which are the residuals conditioned on all the data $t = 1$ to T . The final part of the report covers ‘innovations’, which are residuals conditioned on the data $t = 1$ to $t - 1$.

The MARSS model can be written: $\mathbf{x}_t = \mathbf{B}\mathbf{x}_{t-1} + \mathbf{u} + \mathbf{w}_t$, $\mathbf{y}_t = \mathbf{Z}\mathbf{x}_t + \mathbf{z} + \mathbf{v}_t$, where \mathbf{w}_t and \mathbf{v}_t are multivariate Gaussian error-terms with variance-covariance matrices \mathbf{Q} and \mathbf{R} respectively. The joint conditional residuals are the \mathbf{w}_t and \mathbf{v}_t conditioned on the observed data, which may be incomplete (missing values). Harvey, Koopman and Penzer (1998) show a recursive algorithm for the smoothation residuals (conditioned on all the data). I show an alternate algorithm to compute these residuals using the conditional variances of the states and the conditional covariance between unobserved data and states. This allows one to compute the variance of un-observed residuals (residuals associated with missing or left-out data), which is needed for leave-one-out cross-validation tests. I show how to modify the Harvey et al. algorithm in the case of missing values and how to modify it to return the non-normalized conditional residuals.

Keywords: Time-series analysis, Kalman filter, residuals, maximum-likelihood, vector autoregressive model, dynamic linear model, parameter estimation, state-space

citation: Holmes, E. E. 2014. Computation of standardized residuals for (MARSS) models. Technical Report. arXiv:1411.0045

*Northwest Fisheries Science Center, NOAA Fisheries, Seattle, WA 98112, eli.holmes@noaa.gov, <http://faculty.washington.edu/eeholmes>

1 Overview

This report discusses the computation of the variance of the conditional model and state residuals for MARSS models of the form:

$$\begin{aligned}\mathbf{x}_t &= \mathbf{B}_t \mathbf{x}_{t-1} + \mathbf{u}_t + \mathbf{w}_t, \text{ where } \mathbf{W}_t \sim \text{MVN}(0, \mathbf{Q}_t) \\ \mathbf{y}_t &= \mathbf{Z}_t \mathbf{x}_t + \mathbf{a}_t + \mathbf{v}_t, \text{ where } \mathbf{V}_t \sim \text{MVN}(0, \mathbf{R}_t) \\ \mathbf{X}_0 &\sim \text{MVN}(\boldsymbol{\xi}, \boldsymbol{\Lambda})\end{aligned}\tag{1}$$

The state and model residuals are respectively

$$\begin{aligned}\mathbf{w}_t &= \mathbf{x}_t - \mathbf{B}_t \mathbf{x}_{t-1} - \mathbf{u}_t \\ \mathbf{v}_t &= \mathbf{y}_t - \mathbf{Z}_t \mathbf{x}_t - \mathbf{a}_t\end{aligned}\tag{2}$$

Given a set of observed data \mathbf{y}_t and states \mathbf{x}_t , the model residuals are $\mathbf{y}_t - (\mathbf{Z}_t \mathbf{x}_t + \mathbf{a}_t) = \mathbf{v}_t$. The model residual is a random variable since \mathbf{y}_t and \mathbf{x}_t are drawn from the joint multivariate distribution of \mathbf{Y}_t and \mathbf{X}_t defined by the MARSS equations (Equation 1). The unconditional¹ variance of the model residuals is

$$\text{var}_{XY_t}[\mathbf{V}_t] = \text{var}_{XY_t}[\mathbf{Y}_t - (\mathbf{Z}_t \mathbf{X}_t + \mathbf{a}_t)] = \mathbf{R}_t\tag{3}$$

based on the distribution of \mathbf{V}_t in Equation 1. var_{XY_t} indicates that the integration is over the joint unconditional distribution of \mathbf{X}_t and \mathbf{Y}_t .

Once we have data, \mathbf{R}_t is not the variance-covariance matrix of our model residuals because our residuals are now conditioned² on a set of observed data. There are two types of conditional model residuals used in MARSS analyses: innovations and smoothations. Innovations are the model residuals at time t using the expected value of \mathbf{X}_t conditioned on the data from 1 to $t - 1$. Smoothations are the model residuals using the expected value of \mathbf{X}_t conditioned on all the data, $t = 1$ to T . Smoothations are used in computing standardized residuals for outlier and structural break detection (Harvey et al., 1998; de Jong and Penzer, 1998; Commandeur and Koopman, 2007).

2 Distribution of MARSS smoothation residuals

This section discusses computation of the variance of the model and state residuals conditioned on all the data from $t = 1$ to T . These MARSS residuals are often used for outlier detection and shock detection, and in this case you only need the distribution of the model residuals for the observed values. However if you wanted to do a leave-one-out cross-validation, you would need to know the distribution of the residuals for data points you left out (treated as unobserved). The equations in this report give you the former and the latter, while the algorithm by Harvey et al. (1998) gives only the former.

2.1 Notation and relations

Throughout, I follow the convention that capital letters are random variables and small letters are a realization from the random variable. This only applies to random variables; parameters are not random variables³. Parameters are shown in Roman font while random variables are bold slanted font. Parameters written as capital letters are matrices, while parameters written in small letters are strictly column matrices.

In this report, the distribution over which the integration is done in an expectation or variance is given by the subscript, e.g. $E_A[f(A)]$ indicates an unconditional expectation over the distribution of A without conditioning on another random variable while $E_{A|b}[f(A)|b]$ would indicate an expectation over the distribution of A conditioned on $B = b$; presumably A and B are not independent otherwise $B = b$ would have no effect on A . $E_{A|b}[f(A)|b]$ is a fixed value, not random. It is the expected value when $B = b$. In contrast, $E_{A|B}[f(A)|B]$ denotes the random variable over all the possible $E_{A|b}[f(A)|b]$ given all the possible b values that B might take. The variance of $E_{A|B}[f(A)|B]$ is the variance of this random variable. The variance of

¹meaning not conditioning on any particular set of observed data but rather taking the expectation across all possible values of \mathbf{y}_t and \mathbf{x}_t .

²'conditioned' means that the probability distribution of the residual has changed. The distribution is now the distribution given that $\mathbf{Y} = \mathbf{y}$, say. Expectations and variances $\text{var}[]$ are integrals over the value that a random variable might take multiplied by the probability of that value. When presenting an 'expectation', the probability distribution is normally implicit but for derivations involving conditional expectations, it is important to be explicit about the distribution that is being integrated over.

³in a frequentist framework

$E_{A|b}[f(A)|b]$ in contrast is 0 since it is a fixed value. We will often be working with the random variables, $E_{A|B}[f(A)|B]$ or $\text{var}_{A|B}[f(A)|B]$, inside an expectation or variance: such as $\text{var}_B[E_{A|B}[f(A)|B]]$.

2.1.1 Law of total variance

The “law of total variance” can be written

$$\text{var}_A[A] = \text{var}_B[E_{A|B}[A|B]] + E_B[\text{var}_{A|B}[A|B]] \quad (4)$$

The subscripts on the inner expectations make it explicit that the expectations are being taken over the conditional distributions. $\text{var}_{A|B}[A|B]$ and $E_{A|B}[A|B]$ are random variables because the B in the conditional is a random variable. We take the expectation or variance with it fixed at one value, b , but B can take other values of b also. Going forward, I will write this more succinctly as

$$\text{var}[A] = \text{var}_B[E[A|B]] + E_B[\text{var}[A|B]] \quad (5)$$

It is important to note that $E_{A|b}[A|b]$ and $\text{var}_{A|b}[A|b]$ are fixed values while $E_{A|B}[A|B]$ and $\text{var}_{A|B}[A|B]$ are random variables. When computing $\text{var}_B[E_{A|B}[A|B]]$, we will typically compute $E_{A|b}[A|b]$ and then compute (or infer) the variance or expectation of that over all possible values of b .

The law of total variance will appear in this report in the following form:

$$\text{var}_{XY_t}[f(\mathbf{Y}_t, \mathbf{X}_t)] = \text{var}_{Y^{(1)}}[E_{XY_t|Y^{(1)}}[f(\mathbf{Y}_t, \mathbf{X}_t)|\mathbf{Y}^{(1)}]] + E_{Y^{(1)}}[\text{var}_{XY_t|Y^{(1)}}[f(\mathbf{Y}_t, \mathbf{X}_t)|\mathbf{Y}^{(1)}]] \quad (6)$$

where $f(\mathbf{Y}_t, \mathbf{X}_t)$ is some function of \mathbf{X}_t and \mathbf{Y}_t and $\mathbf{Y}^{(1)}$ is the observed data from $t = 1$ to T .

2.2 Model residuals conditioned on all the data

Define the smoothations $\hat{\mathbf{v}}_t$ as:

$$\hat{\mathbf{v}}_t = \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^T - \mathbf{a}_t, \quad (7)$$

where $\tilde{\mathbf{x}}_t^T$ is $E[\mathbf{X}_t|\mathbf{y}^{(1)}]$. The smoothation is different from \mathbf{v}_t because it uses $\tilde{\mathbf{x}}_t^T$ not \mathbf{x}_t ; \mathbf{x}_t is not known, $\tilde{\mathbf{x}}_t^T$ is its estimate. $\tilde{\mathbf{x}}_t^T$ is output by the Kalman smoother. $\mathbf{y}^{(1)}$ means all the observed data from $t = 1$ to T . $\mathbf{y}^{(1)}$ is a sample from the random variable $\mathbf{Y}^{(1)}$. The unobserved \mathbf{y} will be termed $\mathbf{y}^{(2)}$ and is a sample from the random variable $\mathbf{Y}^{(2)}$. When \mathbf{Y} appears without a superscript, it means both $\mathbf{Y}^{(1)}$ and $\mathbf{Y}^{(2)}$ together. Similarly \mathbf{y} means both $\mathbf{y}^{(1)}$ and $\mathbf{y}^{(2)}$ together—the observed data that we use to estimate $\tilde{\mathbf{x}}_t^T$ and the unobserved data that we do not use and may or may not know. $\hat{\mathbf{v}}_t$ exists for both $\mathbf{y}^{(1)}$ and $\mathbf{y}^{(2)}$, though we might not know $\mathbf{y}^{(2)}$ and thus might not know its corresponding $\hat{\mathbf{v}}_t$. In some cases, however, we do know $\mathbf{y}^{(2)}$; they are data that we left out of our model fitting, in say a k-fold or leave-one-out cross-validation.

$\hat{\mathbf{v}}_t$ is sample from the random variable $\hat{\mathbf{V}}_t$ since \mathbf{Y} is a random variable. We want to compute the mean and variance of this random variable $\hat{\mathbf{V}}_t$ over all possible values that \mathbf{X} and \mathbf{Y} might take. The mean of $\hat{\mathbf{V}}_t$ is 0 and we are concerned only with computing the variance:

$$\text{var}[\hat{\mathbf{V}}_t] = \text{var}_{XY}[\mathbf{Y}_t - \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{Y}^{(1)}] - \mathbf{a}_t] \quad (8)$$

Notice we have an unconditional variance over XY (i.e., over all possible values that \mathbf{X} and \mathbf{Y} can take) on the outside and a conditional expectation over a specific value of $\mathbf{Y}^{(1)}$ on the inside (in the $E[\]$).

From the law of total variance (Equation 4), we can write the variance of the model residuals as

$$\text{var}[\hat{\mathbf{V}}_t] = \text{var}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]] + E_{Y^{(1)}}[\text{var}[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]] \quad (9)$$

2.2.1 First term on right hand side of Equation 9

Notice that the random variable inside the $\text{var}[\]$ in the first term is

$$E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}] = E[(\mathbf{Y}_t + \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{Y}^{(1)}] + \mathbf{a}_t)|\mathbf{Y}^{(1)}] \quad (10)$$

Let’s consider this for a specific value $\mathbf{Y}^{(1)} = \mathbf{y}^{(1)}$.

$$E[\hat{\mathbf{V}}_t|\mathbf{y}^{(1)}] = E[(\mathbf{Y}_t + \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{y}^{(1)}] + \mathbf{a}_t)|\mathbf{y}^{(1)}] = E[\mathbf{Y}_t|\mathbf{y}^{(1)}] + \mathbf{Z}_t E[E_{X_t|y^{(1)}}[\mathbf{X}_t|\mathbf{y}^{(1)}]|\mathbf{y}^{(1)}] + E[\mathbf{a}_t|\mathbf{y}^{(1)}] \quad (11)$$

$E[\mathbf{X}_t|\mathbf{y}^{(1)}]$ is a fixed value, and the expected value of a fixed value is itself. So $E[E[\mathbf{X}_t|\mathbf{y}^{(1)}]|\mathbf{y}^{(1)}] = E[\mathbf{X}_t|\mathbf{y}^{(1)}]$. Thus,

$$E[\hat{\mathbf{V}}_t|\mathbf{y}^{(1)}] = E[\mathbf{Y}_t|\mathbf{y}^{(1)}] + \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{y}^{(1)}] + E[\mathbf{a}_t|\mathbf{y}^{(1)}] \quad (12)$$

We can move the conditional out and write

$$E[\hat{\mathbf{V}}_t|\mathbf{y}^{(1)}] = E[(\mathbf{Y}_t + \mathbf{Z}_t\mathbf{X}_t + \mathbf{a}_t)|\mathbf{y}^{(1)}] = E[\mathbf{V}_t|\mathbf{y}^{(1)}]. \quad (13)$$

The right side is $E[\mathbf{V}_t|\mathbf{y}^{(1)}]$, no ‘hat’ on the \mathbf{V}_t , and this applies for all $\mathbf{y}^{(1)}$. This means that the first term in Equation 9 can be written with no hat on \mathbf{V} :

$$\text{var}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]] = \text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] \quad (14)$$

Using the law of total variance, we can re-write $\text{var}[\mathbf{V}_t]$ as:

$$\text{var}[\mathbf{V}_t] = \text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] + E_{Y^{(1)}}[\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]]. \quad (15)$$

From Equation 15, we can solve for $\text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]]$:

$$\text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] = \text{var}[\mathbf{V}_t] - E_{Y^{(1)}}[\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]] \quad (16)$$

From Equation 3, we know that $\text{var}[\mathbf{V}_t] = \mathbf{R}_t$ (this is the unconditional variance). Thus,

$$\text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] = \mathbf{R}_t - E_{Y^{(1)}}[\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]] \quad (17)$$

The second term in Equation 17 to the right of the equal sign and inside the expectation is $\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]$. This is the variance of \mathbf{V}_t with $\mathbf{Y}^{(1)}$ held at a specific fixed $\mathbf{y}^{(1)}$. The variability in $\text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}]$ (notice $\mathbf{y}^{(1)}$ not $\mathbf{Y}^{(1)}$ now) comes from \mathbf{X}_t (and the rest of the \mathbf{X}) and $\mathbf{Y}^{(2)}$ which are random variables. Let’s compute this variance for a specific $\mathbf{y}^{(1)}$ value.

$$\text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}] = \text{var}[\mathbf{Y}_t - \mathbf{Z}_t\mathbf{X}_t - \mathbf{a}_t|\mathbf{y}^{(1)}]. \quad (18)$$

Notice that there is no E (expectation) on the \mathbf{X}_t ; this is \mathbf{V}_t not $\hat{\mathbf{V}}_t$. \mathbf{a}_t is a fixed value and can be dropped; the \mathbf{Y}_t and \mathbf{X}_t are random variables.

Equation 18 can be written as⁴:

$$\begin{aligned} \text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}] &= \text{var}[\mathbf{Y}_t - \mathbf{Z}_t\mathbf{X}_t|\mathbf{y}^{(1)}] \\ &= \text{var}[-\mathbf{Z}_t\mathbf{X}_t|\mathbf{y}^{(1)}] + \text{var}[\mathbf{Y}_t|\mathbf{y}^{(1)}] + \text{cov}[\mathbf{Y}_t, -\mathbf{Z}_t\mathbf{X}_t|\mathbf{y}^{(1)}] + \text{cov}[-\mathbf{Z}_t\mathbf{X}_t, \mathbf{Y}_t|\mathbf{y}^{(1)}] \\ &= \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top + \tilde{\mathbf{U}}_t^T - \tilde{\mathbf{S}}_t^T\mathbf{Z}_t^\top - \mathbf{Z}_t(\tilde{\mathbf{S}}_t^T)^\top \end{aligned} \quad (19)$$

$\tilde{\mathbf{V}}_t^T = \text{var}[\mathbf{X}_t|\mathbf{y}^{(1)}]$ and is output by the Kalman smoother. $\tilde{\mathbf{U}}_t^T = \text{var}[\mathbf{Y}_t|\mathbf{y}^{(1)}]$ and $\tilde{\mathbf{S}}_t^T = \text{cov}[\mathbf{Y}_t, \mathbf{X}_t|\mathbf{y}^{(1)}]$. The equations for these are given in Holmes (2012) and are output by the `MARSShatyt` function in the MARSS R package⁵. If there were no missing data, i.e. if $\mathbf{y}^{(1)} = \mathbf{y}$, then $\tilde{\mathbf{U}}_t^T$ and $\tilde{\mathbf{S}}_t^T$ would be zero and this would reduce to $\mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top$. But we are concerned with the case where there are missing values. Those missing values need not be for all t . That is, there may be some observed y at time t and some missing y . \mathbf{y}_t is multivariate.

From Equation 19, we know $\text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}]$ for a specific $\mathbf{y}^{(1)}$. We want $E_{Y^{(1)}}[\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]]$ which is its expected value over all possible values of $\mathbf{y}^{(1)}$. $\tilde{\mathbf{V}}_t^T$, $\tilde{\mathbf{U}}_t^T$ and $\tilde{\mathbf{S}}_t^T$ are multivariate normal random variables. The conditional variance of a multivariate normal does not depend on the value that you are conditioning on⁶. This means that $\tilde{\mathbf{V}}_t^T$, $\tilde{\mathbf{U}}_t^T$ and $\tilde{\mathbf{S}}_t^T$ do not depend on $\mathbf{y}^{(1)}$. They only depend on the MARSS model parameters.

⁴ $\text{var}(A + B) = \text{var}(A) + \text{var}(B) + \text{cov}(A, B) + \text{cov}(B, A)$

⁵ $\tilde{\mathbf{U}}_t^T$ is `0tT - tcrossprod(ytT)` in the `MARSShatyt()` output.

⁶ Let the \mathbf{A} be a N-dimensional multivariate normal random variable partitioned into \mathbf{A}_1 and \mathbf{A}_2 with variance-covariance matrix $\Sigma = \begin{bmatrix} \Sigma_1 & \Sigma_{12} \\ \Sigma_{21} & \Sigma_2 \end{bmatrix}$. The variance-covariance matrix of \mathbf{A} conditioned on $\mathbf{A}_1 = \mathbf{a}$ is $\Sigma = \begin{bmatrix} 0 & 0 \\ 0 & \Sigma_2 - \Sigma_{12}\Sigma_1^{-1}\Sigma_{21} \end{bmatrix}$. Notice that \mathbf{a} does not appear in the conditional variance matrix.

Because $\tilde{\mathbf{V}}_t^T$, $\tilde{\mathbf{U}}_t^T$ and $\tilde{\mathbf{S}}_t^T$ only depend on the MARSS parameters values, \mathbf{Q} , \mathbf{B} , \mathbf{R} , etc., the second term in Equation 16, $E_{Y^{(1)}}[\text{var}[\mathbf{V}_t|\mathbf{Y}^{(1)}]]$, is equal to $\text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}]$ (Equation 19). Putting this into Equation 17, we have

$$\text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] = \mathbf{R}_t - \text{var}[\mathbf{V}_t|\mathbf{y}^{(1)}] = \mathbf{R}_t - \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top - \tilde{\mathbf{U}}_t^T + \tilde{\mathbf{S}}_t^T\mathbf{Z}_t^\top + \mathbf{Z}_t(\tilde{\mathbf{S}}_t^T)^\top. \quad (20)$$

Since $\text{var}_{Y^{(1)}}[E[\mathbf{V}_t|\mathbf{Y}^{(1)}]] = \text{var}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]]$ (Equation 14), this means that the first term in Equation 9 is

$$\text{var}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]] = \mathbf{R}_t - \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top - \tilde{\mathbf{U}}_t^T + \tilde{\mathbf{S}}_t^T\mathbf{Z}_t^\top + \mathbf{Z}_t(\tilde{\mathbf{S}}_t^T)^\top. \quad (21)$$

2.2.2 Second term on right hand side of Equation 9

Consider the second term in Equation 9. This term is

$$E_{Y^{(1)}}[\text{var}[(\mathbf{Y}_t - \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{Y}^{(1)}] - \mathbf{a}_t)|\mathbf{Y}^{(1)}]] \quad (22)$$

The middle term is a conditional expectation inside a conditional variance inside an unconditional expectation:

$$E_{Y^{(1)}}[\text{var}[E[\mathbf{X}_t|\mathbf{Y}^{(1)}]|\mathbf{Y}^{(1)}]] \quad (23)$$

Let's solve the inner part for a specific $\mathbf{Y}^{(1)} = \mathbf{y}^{(1)}$. $E[\mathbf{X}_t|\mathbf{y}^{(1)}]$ is a fixed value. Thus $\text{var}[E[\mathbf{X}_t|\mathbf{y}^{(1)}]|\mathbf{y}^{(1)}] = 0$ since the variance of a fixed value is 0. This is true for all $\mathbf{y}^{(1)}$ so the middle term reduces to 0. \mathbf{a}_t is also fixed and its variance is also 0. Thus Equation 22 reduces to $E_{Y^{(1)}}[\text{var}[\mathbf{Y}_t|\mathbf{Y}^{(1)}]] = \tilde{\mathbf{U}}_t^T$. As noted in the previous section, $\tilde{\mathbf{U}}_t^T$ is only a function of the MARSS parameters; it is not a function of $\mathbf{y}^{(1)}$. Thus the second term in Equation 9 is simply $\tilde{\mathbf{U}}_t^T$:

$$E_{Y^{(1)}}[\text{var}[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}]] = \tilde{\mathbf{U}}_t^T \quad (24)$$

2.2.3 Putting together the first and second terms

We can now put the first and second terms in Equation 9 together (Equations 21 and 24) and write out the equation for the variance of the model residuals:

$$\begin{aligned} \text{var}[\hat{\mathbf{V}}_t] &= \mathbf{R}_t - \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top - \tilde{\mathbf{U}}_t^T + \tilde{\mathbf{S}}_t^T\mathbf{Z}_t^\top + \mathbf{Z}_t(\tilde{\mathbf{S}}_t^T)^\top + \tilde{\mathbf{U}}_t^T \\ &= \mathbf{R}_t - \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top + \tilde{\mathbf{S}}_t^T\mathbf{Z}_t^\top + \mathbf{Z}_t(\tilde{\mathbf{S}}_t^T)^\top \end{aligned} \quad (25)$$

This will reduce to $\mathbf{R}_t - \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top$ if \mathbf{y}_t has no missing values and to $\mathbf{R}_t + \mathbf{Z}_t\tilde{\mathbf{V}}_t^T\mathbf{Z}_t^\top$ if \mathbf{y}_t is all missing values. The behavior if \mathbf{y}_t has some missing and some not missing values depends on whether \mathbf{R}_t is a diagonal matrix or not (i.e. if the $\mathbf{y}_t^{(1)}$ and $\mathbf{y}_t^{(2)}$ are correlated).

2.3 State residuals conditioned on the data

The state residuals are $\mathbf{x}_t - (\mathbf{B}_t\mathbf{x}_{t-1} + \mathbf{u}_t) = \mathbf{w}_t$. The unconditional expected value of the state residuals is $E[\mathbf{X}_t - (\mathbf{B}_t\mathbf{X}_{t-1} + \mathbf{u}_t)] = E[\mathbf{W}_t] = 0$ and the unconditional variance of the state residuals is

$$\text{var}[\mathbf{X}_t - (\mathbf{B}_t\mathbf{X}_{t-1} + \mathbf{u}_t)] = \text{var}[\mathbf{W}_t] = \mathbf{Q}_t \quad (26)$$

based on the definition of \mathbf{W}_t . The conditional state residuals (conditioned on the data from $t = 1$ to $t = T$) are defined as

$$\hat{\mathbf{w}}_t = \hat{\mathbf{x}}_t^T - \mathbf{B}_t\hat{\mathbf{x}}_{t-1}^T - \mathbf{u}_t. \quad (27)$$

where $\hat{\mathbf{x}}_t^T = E[\mathbf{X}_t|\mathbf{y}^{(1)}]$ and $\hat{\mathbf{x}}_{t-1}^T = E[\mathbf{X}_{t-1}|\mathbf{y}^{(1)}]$. $\hat{\mathbf{w}}_t$ is a sample from the random variable $\hat{\mathbf{W}}_t$; random over different possible data sets. The expected value of $\hat{\mathbf{W}}_t$ is 0, and we are concerned with computing its variance.

We can write the variance of \mathbf{W}_t (no hat) using the law of total variance.

$$\text{var}[\mathbf{W}_t] = \text{var}_{Y^{(1)}}[E[\mathbf{W}_t|\mathbf{Y}^{(1)}]] + E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]] \quad (28)$$

Notice that $\text{var}_{Y^{(1)}}[E[\mathbf{W}_t|\mathbf{Y}^{(1)}]] = \text{var}[\hat{\mathbf{W}}_t]$.

$$E[\mathbf{W}_t|\mathbf{y}^{(1)}] = E[(\mathbf{X}_t - \mathbf{B}_t\mathbf{X}_{t-1} - \mathbf{u}_t)|\mathbf{y}^{(1)}] = \tilde{\mathbf{x}}_t^T - \mathbf{B}_t\tilde{\mathbf{x}}_{t-1}^T - \mathbf{u}_t = E[\hat{\mathbf{W}}_t|\mathbf{y}^{(1)}] = \hat{\mathbf{w}}_t \quad (29)$$

Thus the random variable $E[\mathbf{W}_t|\mathbf{Y}^{(1)}]$ is $\hat{\mathbf{W}}_t$, and Equation 28 becomes

$$\text{var}[\mathbf{W}_t] = \text{var}[\hat{\mathbf{W}}_t] + E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]] \quad (30)$$

Solve for $\text{var}[\hat{\mathbf{W}}_t]$:

$$\text{var}[\hat{\mathbf{W}}_t] = \text{var}[\mathbf{W}_t] - E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]] \quad (31)$$

The variance in the expectation on the far right for a specific $\mathbf{Y}^{(1)} = \mathbf{y}^{(1)}$ is

$$\text{var}[\mathbf{W}_t|\mathbf{y}^{(1)}] = \text{var}[(\mathbf{X}_t - \mathbf{B}_t\mathbf{X}_{t-1} - \mathbf{u}_t)|\mathbf{y}^{(1)}] \quad (32)$$

\mathbf{u}_t is not a random variable and can be dropped. Thus,

$$\begin{aligned} \text{var}[\mathbf{W}_t|\mathbf{y}^{(1)}] &= \text{var}[(\mathbf{X}_t - \mathbf{B}_t\mathbf{X}_{t-1})|\mathbf{y}^{(1)}] \\ &= \text{var}[\mathbf{X}_t|\mathbf{y}^{(1)}] + \text{var}[\mathbf{B}_t\mathbf{X}_{t-1}|\mathbf{y}^{(1)}] + \text{cov}[\mathbf{X}_t, -\mathbf{B}_t\mathbf{X}_{t-1}|\mathbf{y}^{(1)}] + \text{cov}[-\mathbf{B}_t\mathbf{X}_{t-1}, \mathbf{X}_t|\mathbf{y}^{(1)}] \\ &= \tilde{\mathbf{V}}_t^T + \mathbf{B}_t\tilde{\mathbf{V}}_{t-1}^T\mathbf{B}_t^\top - \tilde{\mathbf{V}}_{t,t-1}^T\mathbf{B}_t^\top - \mathbf{B}_t\tilde{\mathbf{V}}_{t-1,t} \end{aligned} \quad (33)$$

Again this is conditional multivariate normal variance, and its value does not depend on the actual value, $\mathbf{y}^{(1)}$ that we are conditioning on. It depends only on the parameters values, \mathbf{Q} , \mathbf{B} , \mathbf{R} , etc., and is the same for all values of $\mathbf{y}^{(1)}$. So $E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]] = \text{var}[\mathbf{W}_t|\mathbf{y}^{(1)}]$, using any value of $\mathbf{y}^{(1)}$. Thus

$$E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]] = \tilde{\mathbf{V}}_t^T + \mathbf{B}_t\tilde{\mathbf{V}}_{t-1}^T\mathbf{B}_t^\top - \tilde{\mathbf{V}}_{t,t-1}^T\mathbf{B}_t^\top - \mathbf{B}_t\tilde{\mathbf{V}}_{t-1,t} \quad (34)$$

Putting $E_{Y^{(1)}}[\text{var}[\mathbf{W}_t|\mathbf{Y}^{(1)}]]$ from Equation 34 and $\text{var}[\mathbf{W}_t] = \mathbf{Q}_t$ into Equation 31, the variance of the conditional state residuals is

$$\text{var}[\hat{\mathbf{W}}_t] = \mathbf{Q}_t - \tilde{\mathbf{V}}_t^T - \mathbf{B}_t\tilde{\mathbf{V}}_{t-1}^T\mathbf{B}_t^\top + \tilde{\mathbf{V}}_{t,t-1}^T\mathbf{B}_t^\top + \mathbf{B}_t\tilde{\mathbf{V}}_{t-1,t} \quad (35)$$

2.4 Covariance of the conditional model and state residuals

The unconditional model and state residuals, \mathbf{V}_t and \mathbf{W}_t , are independent by definition⁷ (in Equation 1), i.e., $\text{cov}[\mathbf{V}_t, \mathbf{W}_t] = 0$. However the conditional model and state residuals, $\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t]$, are not independent since both depend on $\mathbf{y}^{(1)}$. Using the law of total covariance, we can write

$$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t] = \text{cov}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}], E[\hat{\mathbf{W}}_t|\mathbf{Y}^{(1)}]] + E_{Y^{(1)}}[\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t|\mathbf{Y}^{(1)}]] \quad (36)$$

For a specific value of $\mathbf{Y}^{(1)} = \mathbf{y}^{(1)}$, the covariance in the second term on the right is $\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t|\mathbf{y}^{(1)}]$. Conditioned on a specific value of $\mathbf{Y}^{(1)}$, $\hat{\mathbf{W}}_t$ is a fixed value. $\hat{\mathbf{w}}_t = \tilde{\mathbf{x}}_t^T - \mathbf{B}_t\tilde{\mathbf{x}}_{t-1}^T - \mathbf{u}_t$ and, conditioned on $\mathbf{y}^{(1)}$, $\tilde{\mathbf{x}}_t^T$ and $\tilde{\mathbf{x}}_{t-1}^T$ are fixed values. \mathbf{u}_t is also fixed, because it is a parameter. $\hat{\mathbf{V}}_t$ is not a fixed value because it has $\mathbf{Y}_t^{(2)}$ and that is a random variable. Thus $\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t|\mathbf{y}^{(1)}]$ is the covariance between a random variable and a fixed variable and that covariance is 0.

This eliminates the second right-side term in Equation 36 and the equation reduces to

$$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t] = \text{cov}_{Y^{(1)}}[E[\hat{\mathbf{V}}_t|\mathbf{Y}^{(1)}], E[\hat{\mathbf{W}}_t|\mathbf{Y}^{(1)}]]. \quad (37)$$

Notice that $E[\hat{\mathbf{W}}_t|\mathbf{y}^{(1)}] = E[\mathbf{W}_t|\mathbf{y}^{(1)}]$ and $E[\hat{\mathbf{V}}_t|\mathbf{y}^{(1)}] = E[\mathbf{V}_t|\mathbf{y}^{(1)}]$ since

$$E[\mathbf{W}_t|\mathbf{y}^{(1)}] = E[\mathbf{X}_t|\mathbf{y}^{(1)}] - \mathbf{B}_t E[\mathbf{X}_{t-1}|\mathbf{y}^{(1)}] - \mathbf{u}_t = \tilde{\mathbf{x}}_t^T - \mathbf{B}_t\tilde{\mathbf{x}}_{t-1}^T - \mathbf{u}_t = \hat{\mathbf{w}}_t = E[\hat{\mathbf{W}}_t|\mathbf{y}^{(1)}] \quad (38)$$

and

$$E[\mathbf{V}_t|\mathbf{y}^{(1)}] = E[\mathbf{Y}_t|\mathbf{y}^{(1)}] - \mathbf{Z}_t E[\mathbf{X}_t|\mathbf{y}^{(1)}] - \mathbf{a}_t = E[\mathbf{Y}_t|\mathbf{y}^{(1)}] - \mathbf{Z}_t\tilde{\mathbf{x}}_t^T - \mathbf{a}_t = E[\hat{\mathbf{V}}_t|\mathbf{y}^{(1)}] \quad (39)$$

⁷This independence is specific to the way I have written the MARSS model. It is possible for the model and state residuals to covary. In the MARSS model written in Harvey et al. (1998) form, they do covary.

Thus the right side of Equation 37 can be written in terms of \mathbf{V}_t and \mathbf{W}_t instead of $\hat{\mathbf{V}}_t$ and $\hat{\mathbf{W}}_t$:

$$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t] = \text{cov}_{Y^{(1)}}[\mathbf{E}[\mathbf{V}_t | \mathbf{Y}^{(1)}], \mathbf{E}[\mathbf{W}_t | \mathbf{Y}^{(1)}]]. \quad (40)$$

Using the law of total covariance, we can write:

$$\text{cov}[\mathbf{V}_t, \mathbf{W}_t] = \mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}]] + \text{cov}_{Y^{(1)}}[\mathbf{E}[\mathbf{V}_t | \mathbf{y}^{(1)}], \mathbf{E}[\mathbf{W}_t | \mathbf{y}^{(1)}]] \quad (41)$$

The unconditional covariance of \mathbf{V}_t and \mathbf{W}_t is 0. Thus the left side of Equation 37 is 0 and we can rearrange the equation as

$$\text{cov}_{Y^{(1)}}[\mathbf{E}[\mathbf{V}_t | \mathbf{y}^{(1)}], \mathbf{E}[\mathbf{W}_t | \mathbf{y}^{(1)}]] = -\mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}]] \quad (42)$$

Combining Equation 40 and 42, we get

$$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t] = -\mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}]] \quad (43)$$

and our problem reduces to solving for the conditional covariance of the model and state residuals (right side of Equation 43).

For a specific $\mathbf{Y}^{(1)} = \mathbf{y}^{(1)}$, the conditional covariance $\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}]$ can be written out as

$$\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}] = \text{cov}[\mathbf{Y}_t - \mathbf{Z}_t \mathbf{X}_t - \mathbf{a}_t, \mathbf{X}_t - \mathbf{B}_t \mathbf{X}_{t-1} - \mathbf{u}_t | \mathbf{y}^{(1)}] \quad (44)$$

\mathbf{a}_t and \mathbf{u}_t are fixed values and can be dropped. Thus

$$\begin{aligned} \text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}] &= \text{cov}[\mathbf{Y}_t - \mathbf{Z}_t \mathbf{X}_t, \mathbf{X}_t - \mathbf{B}_t \mathbf{X}_{t-1} | \mathbf{y}^{(1)}] \\ &= \text{cov}[\mathbf{Y}_t, \mathbf{X}_t | \mathbf{y}^{(1)}] + \text{cov}[\mathbf{Y}_t, -\mathbf{B}_t \mathbf{X}_{t-1} | \mathbf{y}^{(1)}] + \text{cov}[-\mathbf{Z}_t \mathbf{X}_t, \mathbf{X}_t | \mathbf{y}^{(1)}] + \text{cov}[-\mathbf{Z}_t \mathbf{X}_t, -\mathbf{B}_t \mathbf{X}_{t-1} | \mathbf{y}^{(1)}] \\ &= \tilde{\mathbf{S}}_t^T - \tilde{\mathbf{S}}_{t,t-1}^T \mathbf{B}_t^\top - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T + \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top \end{aligned} \quad (45)$$

where $\tilde{\mathbf{S}}_t^T = \text{cov}[\mathbf{Y}_t, \mathbf{X}_t | \mathbf{y}^{(1)}]$ and $\tilde{\mathbf{S}}_{t,t-1}^T = \text{cov}[\mathbf{Y}_t, \mathbf{X}_{t-1} | \mathbf{y}^{(1)}]$; the equations for $\tilde{\mathbf{S}}_t^T$ and $\tilde{\mathbf{S}}_{t,t-1}^T$ are given in Holmes (2012) and are output by the MARSShatyt function in the MARSS R package.

$\tilde{\mathbf{V}}_t^T$, $\tilde{\mathbf{V}}_{t,t-1}^T$, $\tilde{\mathbf{S}}_t^T$ and $\tilde{\mathbf{S}}_{t,t-1}^T$ are conditional multivariate normal and are only functions of the MARSS parameters not of $\mathbf{y}^{(1)}$. Thus

$$\mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{Y}^{(1)}]] = \text{cov}[\mathbf{V}_t, \mathbf{W}_t | \mathbf{y}^{(1)}] = \tilde{\mathbf{S}}_t^T - \tilde{\mathbf{S}}_{t,t-1}^T \mathbf{B}_t^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \quad (46)$$

$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t]$ is the negative of this (Equation 43), thus

$$\text{cov}[\hat{\mathbf{V}}_t, \hat{\mathbf{W}}_t] = -\tilde{\mathbf{S}}_t^T + \tilde{\mathbf{S}}_{t,t-1}^T \mathbf{B}_t^\top - \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \quad (47)$$

The Harvey et al. algorithm (shown below) gives the joint distribution of the model residuals at time t and state residuals at time $t+1$. Using the law of total covariance as above, the covariance in this case is

$$\text{cov}_{Y^{(1)}}[\mathbf{E}[\mathbf{V}_t | \mathbf{Y}^{(1)}], \mathbf{E}[\mathbf{W}_{t+1} | \mathbf{Y}^{(1)}]] = -\mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_{t+1} | \mathbf{Y}^{(1)}]] \quad (48)$$

and

$$\begin{aligned} \text{cov}[\mathbf{V}_t, \mathbf{W}_{t+1} | \mathbf{y}^{(1)}] &= \text{cov}[\mathbf{Y}_t - \mathbf{Z}_t \mathbf{X}_t - \mathbf{a}_t, \mathbf{X}_{t+1} - \mathbf{B}_{t+1} \mathbf{X}_t - \mathbf{u}_{t+1} | \mathbf{y}^{(1)}] \\ &= \text{cov}[\mathbf{Y}_t - \mathbf{Z}_t \mathbf{X}_t, \mathbf{X}_{t+1} - \mathbf{B}_{t+1} \mathbf{X}_t | \mathbf{y}^{(1)}] \\ &= \tilde{\mathbf{S}}_{t,t+1}^T - \tilde{\mathbf{S}}_t^T \mathbf{B}_{t+1}^\top - \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t+1}^T + \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{B}_{t+1}^\top \end{aligned} \quad (49)$$

Thus,

$$\text{cov}_{Y^{(1)}}[\mathbf{E}[\mathbf{V}_t | \mathbf{Y}^{(1)}], \mathbf{E}[\mathbf{W}_{t+1} | \mathbf{Y}^{(1)}]] = -\mathbf{E}_{Y^{(1)}}[\text{cov}[\mathbf{V}_t, \mathbf{W}_{t+1} | \mathbf{Y}^{(1)}]] = -\tilde{\mathbf{S}}_{t,t+1}^T + \tilde{\mathbf{S}}_t^T \mathbf{B}_{t+1}^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t+1}^T - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{B}_{t+1}^\top. \quad (50)$$

2.5 Joint distribution of the conditional residuals

We now can write the variance of the joint distribution of the conditional residuals. Define

$$\hat{\epsilon}_t = \begin{bmatrix} \hat{\mathbf{v}}_t \\ \hat{\mathbf{w}}_t \end{bmatrix} = \begin{bmatrix} \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^T - \mathbf{a}_t \\ \tilde{\mathbf{x}}_t^T - \mathbf{B}_t \tilde{\mathbf{x}}_{t-1}^T - \mathbf{u}_t \end{bmatrix}. \quad (51)$$

$\hat{\epsilon}_t$ is a sample drawn from the distribution of the random variable $\hat{\epsilon}_t$. The expected value of $\hat{\epsilon}_t$ over all possible \mathbf{y} is 0 and the variance of $\hat{\epsilon}_t$ is

$$\left[\frac{\text{var}[\hat{V}_t]}{(\text{cov}[\hat{V}_t, \hat{W}_t])^\top} \middle| \frac{\text{cov}[\hat{V}_t, \hat{W}_t]}{\text{var}[\hat{W}_t]} \right] \quad (52)$$

which is

$$\left[\frac{\mathbf{R}_t - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{Z}_t^\top + \tilde{\mathbf{S}}_t^T \mathbf{Z}_t^\top + \mathbf{Z}_t (\tilde{\mathbf{S}}_t^T)^\top}{(-\tilde{\mathbf{S}}_t^T + \tilde{\mathbf{S}}_{t,t-1}^T \mathbf{B}_t^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_t^T - \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top)^\top} \middle| \frac{-\tilde{\mathbf{S}}_t^T + \tilde{\mathbf{S}}_{t,t-1}^T \mathbf{B}_t^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_t^T - \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top}{\mathbf{Q}_t - \tilde{\mathbf{V}}_t^T - \mathbf{B}_t \tilde{\mathbf{V}}_{t-1}^T \mathbf{B}_t^\top + \tilde{\mathbf{V}}_{t,t-1}^T \mathbf{B}_t^\top + \mathbf{B}_t \tilde{\mathbf{V}}_{t-1,t}^\top} \right] \quad (53)$$

If the residuals are defined as in Harvey et al. (1998),

$$\hat{\epsilon}_t = \begin{bmatrix} \hat{\mathbf{v}}_t \\ \hat{\mathbf{w}}_{t+1} \end{bmatrix} = \begin{bmatrix} \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^T - \mathbf{a}_t \\ \tilde{\mathbf{x}}_{t+1} - \mathbf{B}_{t+1} \tilde{\mathbf{x}}_t^T - \mathbf{u}_{t+1} \end{bmatrix} \quad (54)$$

and the variance of $\hat{\epsilon}_t$ is

$$\left[\frac{\text{var}[\hat{V}_t]}{(\text{cov}[\hat{V}_t, \hat{W}_{t+1}])^\top} \middle| \frac{\text{cov}[\hat{V}_t, \hat{W}_{t+1}]}{\text{var}[\hat{W}_{t+1}]} \right] \quad (55)$$

which is

$$\left[\frac{\mathbf{R}_t - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{Z}_t^\top + \tilde{\mathbf{S}}_t^T \mathbf{Z}_t^\top + \mathbf{Z}_t (\tilde{\mathbf{S}}_t^T)^\top}{(-\tilde{\mathbf{S}}_{t,t+1}^T + \tilde{\mathbf{S}}_t^T \mathbf{B}_{t+1}^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t+1}^T - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{B}_{t+1}^\top)^\top} \middle| \frac{-\tilde{\mathbf{S}}_{t,t+1}^T + \tilde{\mathbf{S}}_t^T \mathbf{B}_{t+1}^\top + \mathbf{Z}_t \tilde{\mathbf{V}}_{t,t+1}^T - \mathbf{Z}_t \tilde{\mathbf{V}}_t^T \mathbf{B}_{t+1}^\top}{\mathbf{Q}_{t+1} - \tilde{\mathbf{V}}_{t+1}^T - \mathbf{B}_{t+1} \tilde{\mathbf{V}}_t^T \mathbf{B}_{t+1}^\top + \tilde{\mathbf{V}}_{t+1,t}^T \mathbf{B}_{t+1}^\top + \mathbf{B}_{t+1} \tilde{\mathbf{V}}_{t,t+1}^\top} \right] \quad (56)$$

The above gives the variance of both ‘observed’ model residuals (the ones associated with $\mathbf{y}^{(1)}$) and the unobserved model residuals (the ones associated with $\mathbf{y}^{(2)}$). When there are no missing values in \mathbf{y}_t , the $\tilde{\mathbf{S}}_t^T$ and $\tilde{\mathbf{S}}_{t,t-1}^T$ terms equal 0 and drop out.

3 Harvey et al. 1998 algorithm for the conditional residuals

Harvey et al. (1998, pgs 112-113) give a recursive algorithm for computing the variance of the conditional residuals when the time-varying MARSS equation is written as:

$$\begin{aligned} \mathbf{x}_{t+1} &= \mathbf{B}_{t+1} \mathbf{x}_t + \mathbf{u}_{t+1} + \mathbf{G}_{t+1} \epsilon_t, \\ \mathbf{y}_t &= \mathbf{Z}_t \mathbf{x}_t + \mathbf{a}_t + \mathbf{H}_t \epsilon_t, \end{aligned} \quad (57)$$

$$\text{where } \epsilon_t \sim \text{MVN}(0, \mathbf{I}_{m+n \times m+n}), \mathbf{G}_t \mathbf{G}_t^\top = \mathbf{Q}_t \text{ and } \mathbf{H}_t \mathbf{H}_t^\top = \mathbf{R}_t$$

\mathbf{G}_t has m rows and $m+n$ columns with the last n columns all 0; \mathbf{H}_t has n rows and $m+n$ columns with the last m columns all zero. The algorithm in Harvey et al. (1998) gives the variance of the ‘normalized’ residuals, the ϵ_t . I have modified their algorithm so it returns the ‘non-normalized’ residuals:

$$\varepsilon_t = \begin{bmatrix} \mathbf{H}_t \epsilon_t \\ \mathbf{G}_{t+1} \epsilon_t \end{bmatrix} = \begin{bmatrix} \mathbf{v}_t \\ \mathbf{w}_{t+1} \end{bmatrix}.$$

The Harvey et al. algorithm is a backwards recursion using the following output from the Kalman filter: the one-step ahead prediction covariance \mathbf{F}_t , the Kalman gain \mathbf{K}_t , and $\tilde{\mathbf{V}}_t^{t-1} = \text{var}[\mathbf{X}_t | \mathbf{y}^{(1,1:t-1)}]$. These are output from `MARSSkfss()` in `Sigma`, `Kt`, and `Vtt1`.

4 Algorithm

Start from $t = T$ and work backwards to $t = 1$. At time T , $r_T = 0_{1 \times m}$ and $N_T = 0_{m \times m}$. \mathbf{B}_{t+1} and \mathbf{Q}_{t+1} can be set to NA or 0. They will not appear in the algorithm at time T since $r_T = 0$ and $N_T = 0$. Note that the \mathbf{w} residual and its associated variance and covariance with \mathbf{v} at time T is NA since this residual would be for \mathbf{x}_T to \mathbf{x}_{T+1} .

$$\begin{aligned} \mathbf{Q}_{t+1}^* &= [\mathbf{Q}_{t+1} \quad 0_{m \times n}], \quad \mathbf{R}_t^* = [0_{n \times m} \quad \mathbf{R}_t^*] \\ \mathbf{F}_t &= \mathbf{Z}_t^* \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^{*\top} + \mathbf{R}_t^*, \quad K_t = \mathbf{B}_{t+1} \mathbf{K}_t = \mathbf{B}_{t+1} \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^{*\top} \mathbf{F}_t^{-1} \\ L_t &= \mathbf{B}_{t+1} - K_t \mathbf{Z}_t^*, \quad J_t = \mathbf{Q}_{t+1}^* - K_t \mathbf{R}_t^*, \quad u_t = \mathbf{F}_t^{-1} \hat{\mathbf{v}}_t - K_t^\top r_t \\ r_{t-1} &= \mathbf{Z}_t^{*\top} u_t + \mathbf{B}_{t+1}^\top r_t, \quad N_{t-1} = \mathbf{Z}_t^{*\top} \mathbf{F}_t^{-1} \mathbf{Z}_t^* + L_t^\top N_t L_t \end{aligned} \quad (58)$$

Bolded terms are the same as in Equation 57 or output by `MARSSkfss()`. Unbolded terms are terms used in Harvey et al. (1998). The * on \mathbf{Z}_t and \mathbf{R}_t , indicates that they are the missing value modified versions discussed in Shumway and Stoffer (2006, section 6.4) and Holmes (2012): the rows of \mathbf{Z}_t corresponding to missing rows of \mathbf{y}_t are set to zero and the (i, j) and (j, i) terms of \mathbf{R}_t corresponding the missing rows of \mathbf{y}_t are set to zero. For the latter, this means if the i -th row of \mathbf{y}_t is missing, then then all the (i, j) and (j, i) terms, including (i, i) are set to 0. It is assumed that a missing values modified inverse of \mathbf{F}_t is used: any 0's on the diagonal of \mathbf{F}_t are replaced with 1, the inverse is taken, and 1s on diagonals is replaced back with 0s.

The residuals (Harvey et al., 1998, eqn 24) are

$$\hat{\varepsilon}_t^* = \begin{bmatrix} \hat{\mathbf{v}}_t \\ \hat{\mathbf{w}}_{t+1} \end{bmatrix} = (\mathbf{R}_t^*)^\top u_t + (\mathbf{Q}_{t+1}^*)^\top r_t \quad (59)$$

with mean of 0 ($E_{Y^{(1)}}(\hat{\varepsilon}_t) = 0$) and variance

$$\Sigma_t^* = \text{var}_{Y^{(1)}}(\hat{\varepsilon}_t) = \mathbf{R}_t^{*\top} \mathbf{F}_t^{-1} \mathbf{R}_t^* + J_t^\top N_t J_t \quad (60)$$

The * signifies that these are the missing values modified $\hat{\varepsilon}_t$ and Σ_t ; see comments above.

4.1 Difference in notation

In Equation 20 in Harvey et al. (1998), their T_t is my \mathbf{B}_{t+1} and their $H_t H_t^\top$ is my \mathbf{Q}_{t+1} . Notice the difference in the time indexing. My time indexing on \mathbf{B} and \mathbf{Q} matches the left \mathbf{x} while in theirs, T and H indexing matches the right \mathbf{x} . Thus in my implementation of their algorithm (Harvey et al., 1998, eqns. 21-24), \mathbf{B}_{t+1} appears in place of T_t and \mathbf{Q}_{t+1} appears in place of H_t . See comments below on normalization and the difference between \mathbf{Q} and H .

Harvey et al. (1998, eqns. 19, 20) use G_t to refer to the $\text{chol}(\mathbf{R}_t)^\top$ (essentially) and H_t to refer to $\text{chol}(\mathbf{Q}_t)^\top$. I have replaced these with \mathbf{R}_t^* and \mathbf{Q}_t^* , respectively, which causes my variant of their algorithm (Equation 58) to give the ‘non-normalized’ variance of the residuals. The residuals function in the MARSS package has an option to give either normalized or non-normalized residuals.

\mathbf{K}_t is the Kalman gain output by the MARSS package. The Kalman gain as used in the Harvey et al. (1998) algorithm is $K_t = \mathbf{B}_{t+1} \mathbf{K}_t$. Notice that Equation 21 in Harvey et al. (1998) has $H_t G_t^\top$ in the equation for K_t (Kalman gain). This is the covariance of the state and observation errors, which is allowed given the way they write the errors in their Equations 19 and 20. The way the MARSS package model is written, the state and observation errors are independent of each other. Thus $H_t G_t^\top = 0$ and this term drops out of the \mathbf{K}_t equation.

4.2 Computing the standardized residuals

The standardized residuals are computed by multiplying $\hat{\varepsilon}_t$ by the inverse of the square root of the variance-covariance matrix from which $\hat{\varepsilon}_t$ is “drawn”:

$$(\Sigma_t^*)^{-1/2} \hat{\varepsilon}_t^* \quad (61)$$

Notice that the missing values modified $\hat{\varepsilon}_t^*$ and Σ_t^* are used. if the i -th row of \mathbf{y}_t is missing, the i -th row of $\hat{\varepsilon}_t$ is set to 0 and the i -th row and column of Σ_t is set to all 0. There will be 0s on the diagonal of Σ_t^* so your code will need to deal with these when taking the inverse.

5 Distribution of the MARSS innovation residuals

One-step-ahead predictions (innovations) are often shown for MARSS models and these are used for likelihood calculations. Innovations are the prediction of \mathbf{y} given data up to $t - 1$. Although not termed 'innovations', we can also compute the predictions of \mathbf{x} conditioned on the data up to $t - 1$. This section gives the residual variance for these one-step-ahead predictions.

5.0.1 Variance of the one-step-ahead model residuals

Define the innovations $\tilde{\mathbf{v}}_t$ as:

$$\tilde{\mathbf{v}}_t = \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^{t-1} - \mathbf{a}_t, \quad (62)$$

where $\tilde{\mathbf{x}}_t^{t-1}$ is the $E[\mathbf{X}_t | \mathbf{y}^{(1),t-1}]$ (expected value of \mathbf{X}_t conditioned on the data up to time $t - 1$). The random variable, innovations over all possible $\mathbf{y}^{(1),t-1}$, is $\tilde{\mathbf{V}}_t$. Its mean is 0 and we want to find its variance.

The derivation of the variance of $\tilde{\mathbf{V}}_t$ follows the same steps as smoothations $\hat{\mathbf{V}}_t$, and we can write the variance as:

$$\text{var}[\tilde{\mathbf{V}}_t] = \mathbf{R}_t - \mathbf{Z}_t \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^\top + \tilde{\mathbf{S}}_t^{t-1} \mathbf{Z}_t^\top + \mathbf{Z}_t (\tilde{\mathbf{S}}_t^{t-1})^\top \quad (63)$$

where the $\tilde{\mathbf{V}}_t^{t-1}$ and $\tilde{\mathbf{S}}_t^{t-1}$ are now conditioned on only the data from 1 to $t - 1$. $\tilde{\mathbf{S}}_t^{t-1} = \text{cov}[\mathbf{Y}_t, \mathbf{X}_t | \mathbf{y}^{(1),t-1}]$ which is $\text{cov}[\mathbf{Z}_t \mathbf{X}_t + \mathbf{a}_t, \mathbf{X}_t | \mathbf{y}^{(1),t-1}] = \mathbf{Z}_t \tilde{\mathbf{V}}_t^{t-1}$ and therefore, $\mathbf{Z}_t (\tilde{\mathbf{S}}_t^{t-1})^\top = \tilde{\mathbf{S}}_t^{t-1} (\mathbf{Z}_t)^\top = \mathbf{Z}_t \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^\top$. Thus Equation 63 reduces to

$$\text{var}[\tilde{\mathbf{V}}_t] = \mathbf{R}_t + \mathbf{Z}_t \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^\top. \quad (64)$$

5.1 State residuals conditioned on the data

Define the state residuals conditioned on the data from 1 to $t - 1$ as $\tilde{\mathbf{w}}_t$.

$$\tilde{\mathbf{w}}_t = \tilde{\mathbf{x}}_t^{t-1} - \mathbf{B}_t \tilde{\mathbf{x}}_{t-1}^{t-1} - \mathbf{u}_t, \quad (65)$$

where $\tilde{\mathbf{x}}_t^{t-1}$ is the $E[\mathbf{X}_t | \mathbf{y}^{(1),t-1}]$ (expected value of \mathbf{X}_t conditioned on the data up to time $t - 1$). However $\tilde{\mathbf{x}}_t^{t-1} = E[\mathbf{X}_t | \mathbf{y}^{(1),t-1}] = \mathbf{B}_t \tilde{\mathbf{x}}_{t-1}^{t-1} + \mathbf{u}_t$. Thus $\tilde{\mathbf{w}}_t = 0$.

5.2 Covariance of the conditional model and state residuals

Since $\tilde{\mathbf{w}}_t = 0$ for all $\mathbf{y}^{(1)}$:

$$\text{cov}[\tilde{\mathbf{V}}_t, \tilde{\mathbf{W}}_t] = 0 \text{ and } \text{cov}[\tilde{\mathbf{V}}_t, \tilde{\mathbf{W}}_{t+1}] = 0 \quad (66)$$

5.3 Joint distribution of the conditional residuals

We now write the variance of the joint distribution of the conditional one-step ahead residuals. Define

$$\tilde{\mathcal{E}}_t = \begin{bmatrix} \tilde{\mathbf{v}}_t \\ \tilde{\mathbf{w}}_t \end{bmatrix} = \begin{bmatrix} \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^{t-1} - \mathbf{a}_t \\ \tilde{\mathbf{x}}_t^{t-1} - \mathbf{B}_t \tilde{\mathbf{x}}_{t-1}^{t-1} - \mathbf{u}_t \end{bmatrix}. \quad (67)$$

where $\tilde{\mathbf{x}}_t^{t-1}$ and $\tilde{\mathbf{x}}_{t-1}^{t-1}$ are conditioned on the observed \mathbf{y} from $t = 1$ to $t - 1$. The expected value of $\tilde{\mathcal{E}}_t$ over all possible \mathbf{y} up to time $t - 1$ is 0 and the variance of $\tilde{\mathcal{E}}_t$ is

$$\left[\frac{\text{var}[\tilde{\mathbf{V}}_t]}{(\text{cov}[\tilde{\mathbf{V}}_t, \tilde{\mathbf{W}}_t])^\top} \middle| \frac{\text{cov}[\tilde{\mathbf{V}}_t, \tilde{\mathbf{W}}_t]}{\text{var}[\tilde{\mathbf{W}}_t]} \right] = \left[\frac{\mathbf{R}_t + \mathbf{Z}_t \tilde{\mathbf{V}}_t^{t-1} \mathbf{Z}_t^\top}{0} \middle| \frac{0}{0} \right] \quad (68)$$

If the residuals are defined as in Harvey et al. (1998),

$$\hat{\mathcal{E}}_t = \begin{bmatrix} \tilde{\mathbf{v}}_t \\ \tilde{\mathbf{w}}_{t+1} \end{bmatrix} = \begin{bmatrix} \mathbf{y}_t - \mathbf{Z}_t \tilde{\mathbf{x}}_t^{t-1} - \mathbf{a}_t \\ \tilde{\mathbf{x}}_{t+1}^{t-1} - \mathbf{B}_{t+1} \tilde{\mathbf{x}}_t^{t-1} - \mathbf{u}_{t+1} \end{bmatrix} \quad (69)$$

$\tilde{\mathbf{W}}_{t+1} = 0$, and the variance-covariance matrix is again Equation 68.

References

- Commandeur, J. J. and Koopman, S. J. (2007). *An introduction to state space time series analysis*. Practical Econometrics. Oxford University Press, Oxford.
- de Jong, P. and Penzer, J. (1998). Diagnosing shocks in time series. *Journal of the American Statistical Association*, 93(442):796–806.
- Harvey, A., Koopman, S. J., and Penzer, J. (1998). Messy time series: a unified approach. *Advances in Econometrics*, 13:103–143.
- Holmes, E. E. (2012). Derivation of the EM algorithm for constrained and unconstrained MARSS models. Technical report, arXiv:1302.3919 [stat.ME].
- Shumway, R. and Stoffer, D. (2006). *Time series analysis and its applications*. Springer-Science+Business Media, LLC, New York, New York, 2nd edition.